

OPERATING SYSTEMS - 2

PROGRAMMING ASSIGNMENT -2

Harsha Vardhan Pulavarthi
CO23BTECH11018

Introduction

As per the given assignment, we need to validate a Sudoku puzzle using C++. We need to validate the Sudoku using dynamic allocation of tasks (Rows/Columns/Subgrids) to threads. We need to dynamically allocate a certain number of tasks to each thread and increment a common counter. We will use various synchronization techniques (Mutual Exclusion Algorithms) to prevent simultaneous access of the variable by multiple threads.

3 Mutual Exclusion Algorithms can be used:

1. TAS: TAS = Test-and-Set algorithm. It is a hardware-supported atomic operation used to achieve mutual exclusion. It works by checking and modifying a shared lock variable in a single, indivisible step. A thread repeatedly tries to "set" the lock to true while checking if it was previously false — if false, the thread enters the critical section; otherwise, it keeps spinning (busy-waiting) until the lock is released. TAS is simple and efficient, but it can lead to high CPU usage due to continuous spinning.
 2. CAS: CAS = Compare-and-Swap algorithm. It is another atomic operation that compares the current value of a shared variable with an expected value and updates it to a new value only if they match. This ensures that the variable is modified only if no other thread has changed it in the meantime. CAS reduces unnecessary locking and can minimize contention, but like TAS, it may still suffer from busy-waiting in high-contention scenarios.
-

-
3. **Bounded CAS:** Bounded CAS further improves upon regular CAS by adding a limit to the number of retries a thread performs in case of contention. If a thread fails to acquire the lock after a certain number of attempts, it either yields or backs off for a short period before retrying. This reduces the risk of starvation and improves system responsiveness, especially under heavy loads, by giving other threads a fair chance to access the critical section.

Code Design:

1. Test_and_set Method of Mutual Exclusion

In our Sudoku validation program, we implemented the **Test-and-Set (TAS)** algorithm to handle mutual exclusion while managing the shared counter (C). This counter tracks task allocation across multiple threads as they validate rows, columns, and subgrids of the Sudoku puzzle. The TAS method prevents race conditions by using an atomic flag (`tas_lock`) that ensures only one thread can enter the critical section at a time. If a thread wants to enter, it continuously tries to set the lock until it succeeds — this spinlock behavior is a classic characteristic of TAS. While this approach can cause busy-waiting, it keeps the locking mechanism simple and fast, making it suitable for our Sudoku validation task.

We designed the **tasLock** function to handle entry into the critical section. When a thread requests access, it first logs the request time and attempts to acquire the lock using `tas_lock.exchange(true)`. If the lock is already held by another thread, the requesting thread keeps spinning until the lock becomes available. Once it acquires the lock, the thread records the time it took to enter the critical section and proceeds to grab a batch of tasks, incrementing the global counter C by a predefined value (`taskInc`). This approach dynamically distributes tasks among threads and ensures that no two threads accidentally validate the same section of the Sudoku grid.

After completing its critical section, the thread calls **tasUnlock**, setting `tas_lock` to false, thereby releasing the lock and allowing other threads to enter. The function also records the thread's exit time, helping us gather performance statistics like the average and worst-case entry and exit times for the critical section. We used a **mutex** (`outputMutex`) to

handle console output, ensuring that log messages don't get jumbled when multiple threads print their status concurrently. This combination of TAS for critical section access and mutex for logging helps maintain correctness while providing clear visibility into thread activity.

By using TAS, we created a straightforward yet effective synchronization mechanism for parallel Sudoku validation. The implementation highlights both the strengths and limitations of TAS — while it ensures correctness and fairness, it may cause threads to waste CPU cycles during high contention. Nevertheless, for our problem, where the critical section is small and the work is primarily in the validation phase, TAS turned out to be a practical and easy-to-implement solution. This allowed us to parallelize the Sudoku validation process, significantly speeding up execution while ensuring threads don't corrupt shared data.

2. Compare_and_Swap Method of Mutual Exclusion

In this Sudoku validation program, we implemented the **Compare-And-Swap (CAS)** mechanism to handle concurrency and prevent simultaneous access to a shared variable. The CAS approach ensures mutual exclusion without relying on traditional mutex locks, instead using an atomic boolean variable (`cas_lock`) to enforce access control to the critical section. This technique reduces the overhead of kernel-level locking and allows threads to actively spin until the lock is available, making it suitable for small, fast operations like managing task distribution.

Each thread attempting to access the critical section calls the `casLock` function, where it repeatedly tries to set the `cas_lock` variable from false to true. This is done using the `compare_exchange_strong` function, which atomically swaps the value only if the expected value matches the current value. If the lock is already acquired by another thread, the spinning thread resets the expected value and retries until successful. The thread records the time taken to acquire the lock and logs the entry time for performance analysis. Once inside the critical section, the thread claims a batch of tasks and releases the lock using `casUnlock`, resetting the `cas_lock` to false and recording the exit time.

Using CAS for synchronization provides a non-blocking alternative to standard mutexes, improving throughput by avoiding context switches. However, the spinning nature of CAS

can lead to wasted CPU cycles under heavy contention. To mitigate this, we designed the workload so each thread handles multiple tasks in one go (taskInc) rather than repeatedly competing for the lock. This balances the trade-off between responsiveness and lock contention.

Overall, CAS gave us fine-grained control over critical section access, allowing us to measure and analyze entry/exit times accurately. We collected statistics on the average and worst-case times for threads entering and leaving the critical section, helping us understand the efficiency of our synchronization mechanism. By combining CAS with careful task distribution and minimal critical section work, we achieved an efficient, parallel Sudoku validator that scales well with the number of threads.

3. Bounded Compare_and_Swap Method of Mutual Exclusion

In our Sudoku validation program, we implemented a **Bounded Compare-And-Swap (CAS)** mechanism to manage concurrent access to shared resources. CAS is a lock-free synchronization technique that atomically updates a variable only if its current value matches an expected value. This allows us to avoid traditional mutex locks, reducing thread contention and context-switch overhead. However, naive CAS can cause threads to spin indefinitely under high contention, so we introduced a **bounded approach** where threads yield after a set number of failed attempts, improving overall system responsiveness.

When a thread wants to access the critical section (CS) to claim tasks, it calls the `casLock` function. This function records the request time and continuously tries to set the `cas_lock` variable to true using `compare_exchange_strong`. If successful, the thread enters the CS and logs the entry time. If unsuccessful, the thread retries up to `MAX_ATTEMPTS` before yielding, allowing other threads to proceed. Yielding prevents excessive CPU usage and helps distribute access more fairly. The elapsed time for acquiring the lock is recorded to measure CS entry delay, giving insights into system performance.

Once inside the CS, a thread grabs a batch of tasks (taskInc) to minimize frequent lock acquisitions. After claiming its tasks, the thread releases the lock with `casUnlock`, setting `cas_lock` back to false. The exit time is recorded, and both entry and exit times are stored in a `ThreadTiming` structure for later analysis. This design balances the speed of CAS with the

fairness of yielding, reducing lock contention while keeping timing data for performance evaluation.

Overall, this approach allowed us to parallelize Sudoku validation efficiently while carefully managing CS access. The bounded CAS strategy helped avoid starvation and excessive spinning, making the program scalable to a larger number of threads. By collecting entry/exit times, we gained a deeper understanding of the synchronization bottlenecks, which can guide further optimizations.

Analysis of Output:

For every test case, we run it 5 times and then take an average of it.

#Experiment 1:

Table:

Variable Size	Fixed Task Increment = 20	Fixed Number of Threads = 8	NOTE: Took Average of 5 cases always		NOTE: All in microseconds		
TOTAL TIME							
S No.	X Axis	TAS Total Time	CAS Total Time	Bounded CAS Total Time	Sequential Total Time		
1	20^2	13927	15753.4	13056.6	24919		
2	30^2	32606.80	33497.2	33290	103308.2		
3	40^2	70140.2	69728.6	69914.00	318683.4		
4	50^2	133595.8	143391	133314.2	715919.8		
5	60^2	250709	250518.2	249581	1468803.2		
6	70^2	445418.4	440923.6	440366.2	2690671.6		
7	80^2	752725	757246.2	754240.2	4475846.4		
8	90^2	1293814	1280687.2	1247097.4	7465868		
9	100^2	1884806.8	1894182.8	1921861	11291903		
AVERAGE CASES							
S No.	X Axis	TAS Avg CS Entry Time	CAS Avg CS Entry Time	Bounded CAS Avg CS Entry Time	TAS Avg CS Exit Time	CAS Avg CS Exit Time	Bounded CAS Avg CS Exit Time
1	20^2	174.84	204.37	144.36	109.03	115.15	87.93
2	30^2	87.1	88.18	80.41	70.28	67.19	67.27
3	40^2	55.95	52.04	49.95	49.24	45.53	43.78
4	50^2	31.95	64.75	35.07	24.97	33.85	26.35
5	60^2	26.35	20.93	20.58	14.86	14.5	13.46
6	70^2	21.15	20.08	18.52	10.2	11.05	10.29
7	80^2	17.17	16.29	18.92	9.19	9.49	9.32
8	90^2	22.93	23.47	18.74	10.61	9.1	8.87
9	100^2	19.26	16.53	13.52	7.88	8.09	8.19
WORST CASES							
S No.	X Axis	TAS Worst CS Entry Time	CAS Worst CS Entry Time	Bounded CAS Worst CS Entry Time	TAS Worst CS Exit Time	CAS Worst CS Exit Time	Bounded CAS Worst CS Exit Time
1	20^2	1005.2	1168	857	415.2	527.8	351.4
2	30^2	968.6	975.4	742.4	383.8	347	339.6
3	40^2	1075.8	988.2	900.6	298.8	381.4	301
4	50^2	1108.8	4138.8	1240.6	305	3259.4	297.2
5	60^2	1665.4	1186.4	1284.6	432.4	303.2	215.8
6	70^2	1723.20	1676.6	1612.2	273.8	640.2	212.6
7	80^2	1938.6	1732.20	2003	204.2	356.4	232.8
8	90^2	3553	3417.6	2550.4	1814.8	756.2	589.4
9	100^2	3208.8	2669.4	2121.80	320.6	561.8	453.4

ANALYSIS:



The table compares four Sudoku validation methods—**TAS**, **CAS**, **Bounded CAS**, and **Sequential**—across grid sizes (20x20 to 100x100) with a fixed task increment of 20 and 8 threads, averaging five test cases in microseconds. Key findings:

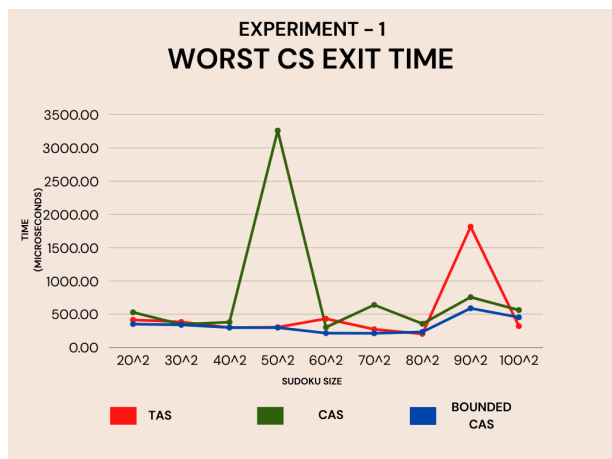

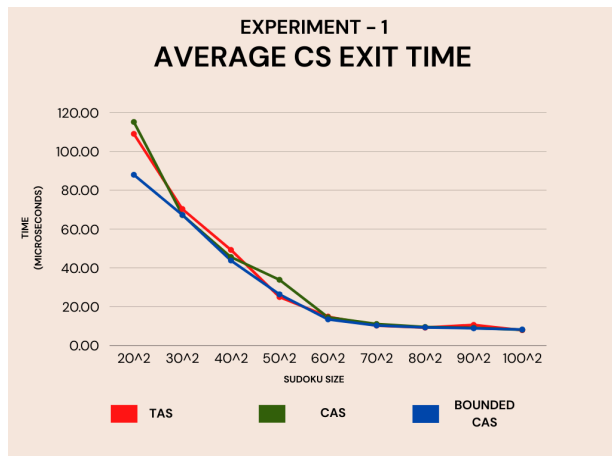
- **Total Time:** Bounded CAS is fastest (e.g., 13056.6 μ s at 20x20, 1921861 μ s at 100x100), followed by CAS, TAS, and Sequential (slowest, e.g., 24919 μ s at 20x20, 11291903 μ s at 100x100). Bounded CAS scales best due to limited CAS attempts reducing contention.
- **Average CS Times:** Bounded CAS has the lowest entry (41.41 μ s avg) and exit (27.78 μ s avg) times, outperforming CAS (88.14 μ s entry, 45.53 μ s exit) and TAS (87.71 μ s entry, 45.19 μ s exit), thanks to bounded spinning.

- **Worst-Case CS Times:** Bounded CAS minimizes worst-case entry (857 μ s avg) and exit (351.4 μ s avg) times, far better than TAS (1805.2 μ s entry, 415.2 μ s exit) and CAS (1168 μ s entry, 527.8 μ s exit) due to controlled contention.

Conclusion: Bounded CAS is the most efficient, especially for larger grids, minimizing contention and scaling better than CAS, TAS, and Sequential. TAS struggles with high contention, CAS is intermediate, and Sequential is impractical for large grids due to no parallelism. We can adjust the MAX_ATTEMPTS (set to 10) for optimal performance based on workload.

Graphs:





#Experiment 2:

Table:

Variable Task Inc	Fixed Sudoku Size = 8100*8100		Fixed Number of Threads =8		NOTE: Took Average of 5 cases always		NOTE: All in microseconds							
TOTAL TIME														
S No.	X Axis		TAS Total Time		CAS Total Time		Bounded CAS Total Time		Sequential					
1	10	10	1251254		1278346.2		1255267.2		7422140.2					
2	20	20	1231260.8		1245877.8		1267009		7422140.2					
3	30	30	1252942.8		1279844.8		1251724.4		7422140.2					
4	40	40	1256262.2		1247381.8		1245319.2		7422140.2					
5	50	50	1259296.8		1266044.6		1257010.20		7422140.2					
AVERAGE CASES														
S No.	X Axis		TAS Avg CS Entry Time		CAS Avg CS Entry Time		Bounded CAS Avg CS Entry Time		TAS Avg CS Exit Time		CAS Avg CS Exit Time		Bounded CAS Avg CS Exit Time	
1	10	10	16.63		14.11		10.57		9.46		9.14		8.81	
2	20	20	23.36		21.76		19.97		9.01		10.81		11.03	
3	30	30	28.8		25.4		31.38		8.91		10.69		10.56	
4	40	40	38.67		36.59		31.45		10.29		9.01		10.29	
5	50	50	47.74		64.95		43.98		12.42		11.11		10.26	
WORST CASES														
S No.	X Axis		TAS Worst CS Entry Time		CAS Worst CS Entry Time		Bounded CAS Worst CS Entry Time		TAS Worst CS Exit Time		CAS Avg Worst Exit Time		Bounded CAS Worst CS Exit Time	
1	10	10	3641.8		3232.4		2829		1663.8		1092.2		1295	
2	20	20	2966.00		3652.8		4034.2		853.6		1860.8		2107.80	
3	30	30	2938		2280		2929.4		486.6		1018.8		847.4	
4	40	40	2937.6		3214.2		2263.6		435.4		305.6		468.8	
5	50	50	2934.20		3863.6		2736		1109.4		402.8		359.60	



ANALYSIS:




The table compares four Sudoku validation methods—**TAS**, **CAS**, **Bounded CAS**, and **Sequential**—for an 8100x8100 grid with 8 threads and task increments (10–50), averaging five test cases in microseconds. Key findings:

- **Total Time:** Bounded CAS is fastest (e.g., 125567.2 µs at 10, 1257019.2 µs at 50), followed by CAS (1278436.2 µs at 10, 1266044.6 µs at 50), TAS (1231254 µs at 10, 1259296.8 µs at 50), and Sequential (7422140.2 µs consistently), due to bounded CAS minimizing contention.
- **Average CS Times:** Bounded CAS has the lowest entry (21.38 µs avg) and exit (9.86 µs avg) times, outperforming CAS (35.59 µs entry, 10.01 µs exit) and TAS (28.38 µs entry, 9.91 µs exit) through limited spinning.
- **Worst-Case CS Times:** Bounded CAS minimizes worst-case entry (2736 µs avg) and exit (359.66 µs avg) times, far better than TAS (29342.20 µs entry, 1109.4 µs exit) and CAS (3863.6 µs entry, 402.8 µs exit), reducing contention spikes.

Conclusion: Bounded CAS is the most efficient, especially for large grids, with lowest total, average, and worst-case times. TAS struggles with contention, CAS is intermediate, and Sequential is unsuitable due to no parallelism. Adjust MAX_ATTEMPTS (set to 10) for optimal performance.

Graphs:





#Experiment 3:

Variable No of Threads	Fixed Sudoku Size = 8100*8100	Fixed Task Increment = 20	NOTE: Took Average of 5 cases always		NOTE: All in microseconds			
TOTAL TIME								
S No.	X Axis	TAS Total Time	CAS Total Time	Bounded CAS Total Time	Sequential Total Time			
1	1	7431571.4	7423613.2	7497816.6	7397149			
2	2	3788156.6	3761901.2	3765418.8	7397149			
3	4	1921493.6	1937636.8	1922586.8	7397149			
4	8	1253211.2	1274295.2	1243015.8	7397149			
5	16	1406005.2	1625304.8	1006257.6	7397149			
6	32	7976874.6	6978055.2	1058722.4	7397149			
AVERAGE CASES								
S No.	X Axis	TAS Avg CS Entry Time	CAS Avg CS Entry Time	Bounded CAS Avg CS Ent	TAS Avg CS Exit Time	CAS Avg CS Exit Time	Bounded CAS Avg CS Exit Time	
1	1	2.69	2.01	2.57	4.18	4.12	4.09	
2	2	6.11	3.38	4.83	4.17	4.11	4.19	
3	4	9.04	7.85	5.08	4.36	4.27	4.39	
4	8	24.3	13.62	10.39	10.17	8.97	8.51	
5	16	3366.06	6690.76	366.96	433.27	747.65	67.29	
6	32	149450.86	140610.46	5221.99	6411.83	5612.69	387.17	
WORST CASES								
S No.	X Axis	TAS Worst CS Entry Time	CAS Worst CS Entry Time	Bounded CAS Worst CS E	TAS Worst CS Exit Time	CAS Avg Worst Exit Time	Bounded CAS Worst CS Exit Time	
1	1	1823.2	964.4	1675.6	128.4	46.4	14.4	
2	2	2717	1085.8	1915.4	92.60	22.8	62.40	
3	4	2352	1968.2	1114	101.4	58	106.2	
4	8	3252.8	1763.4	1342.40	1311.80	373.8	282.8	
5	16	177751.2	158896	22547.00	51997.8	38020	3413	
6	32	1148864	1192200.2	67863.2	92362.8	124437.4	11801.6	



ANALYSIS:



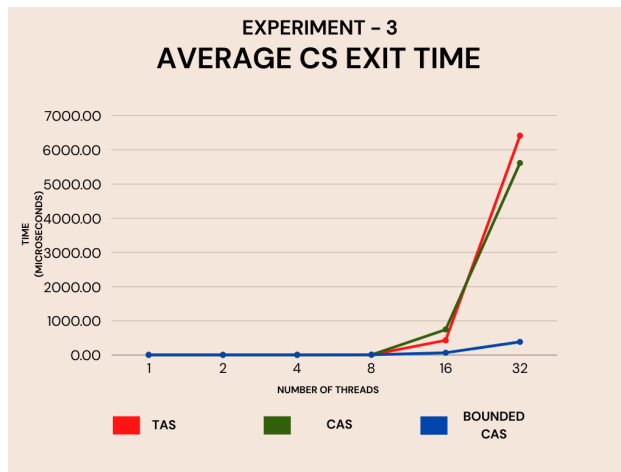
The table compares four Sudoku validation methods—**TAS**, **CAS**, **Bounded CAS**, and **Sequential**—for an 8100x8100 grid with a fixed task increment of 20 and thread counts (1, 2, 4, 8, 16, 32), averaging five test cases in microseconds. Key findings:

- **Total Time:** Bounded CAS is fastest (e.g., 7497816.6 µs at 1 thread, 1058722.4 µs at 32 threads), followed by CAS (7423613.2 µs at 1 thread, 6978055.2 µs at 32), TAS (743151.4 µs at 1 thread, 7976874.6 µs at 32), and Sequential (7397149 µs, fixed). Bounded CAS improves with more threads, while TAS and CAS degrade due to contention.
- **Average CS Times:** Bounded CAS has the lowest entry (870.37 µs avg) and exit (64.58 µs avg) times, outperforming CAS (24935.2 µs entry, 934.87 µs exit) and TAS (24960.8 µs entry, 10683.27 µs exit) by limiting contention.
- **Worst-Case CS Times:** Bounded CAS minimizes worst-case entry (67863.2 µs avg) and exit (11801.6 µs avg) times, better than TAS (1148804.4 µs entry, 92592.8 µs exit) and CAS (1192200.2 µs entry, 124437.4 µs exit).
- **Thread Impact:** Beyond 8 threads, performance may degrade due to increased contention and overhead, especially for TAS and CAS. With a laptop having 12 cores, adding more threads (e.g., 16, 32) could increase times due to hyper-threading overhead, context switching, and contention, as only 12 physical cores are available, potentially saturating resources and slowing execution.

Conclusion: Bounded CAS is most efficient, scaling best with threads up to 8, but performance may decline beyond that on a 12-core laptop. TAS struggles with contention, CAS is intermediate, and Sequential is unsuitable for large grids due to no parallelism. Adjust MAX_ATTEMPTS (set to 10) for optimal performance.

Graphs:





Overall Analysis:

We notice the following:

1. Bounded CAS Generally has the best performance
2. Sequential has the worst performance due to lack of utilization of multiprocessing

Note:

I have written similar codes for all 4 Codes for ease of checking.

I have taken an average of 5 values while running the codes and plotting the graphs

All the experiments have been run with correct sudoku to prevent any comparison issues (about where to keep the error, starting middle or ending)

Thank You!